

ADS Assignment 3

Problem Statement: House Price Prediction

Description:-

House price prediction is a common problem in the real estate industry and involves predicting the selling price of a house based on various features and attributes. The problem is typically approached as a regression problem, where the target variable is the price of the house, and the features are various attributes of the house. The features used in house price prediction can include both quantitative and categorical variables, such as the number of bedrooms, house area, bedrooms, furnished, nearness to main road, and various amenities such as a garage and other factors that may influence the value of the property. Accurate predictions can help agents and appraisers price homes correctly, while homeowners can use the predictions to set a reasonable asking price for their properties. Accurate house price prediction can also be useful for buyers who are looking to make informed decisions about purchasing a property and obtaining a fair price for their investment.

Attribute Information:

Name - Description

- 1- Price-Prices of the houses
- 2- Area- Area of the houses
- 3- Bedrooms- No of house bedrooms
- 4- Bathrooms- No of bathrooms
- 5- Stories- No of house stories
- 6- Main Road- Weather connected to Main road
- 7- Guestroom-Weather has a guest room
- 8- Basement-Weather has a basement
- 9- Hot water heating- Weather has a hot water heater
- 10-Airconditioning-Weather has a air conditioner
- 11-Parking- No of house parking
- 12-Furnishing Status-Furnishing status of house

1. Download the dataset: titanic.csv

2. Load the dataset.

```
In [1]: # Loading necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # Loading the dataset
df = pd.read_csv('Housing.csv')
df.head()
```

```
Out[2]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterhe
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	

```
In [3]: df.shape
```

```
Out[3]: (545, 12)
```


```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 545 entries, 0 to 544  
Data columns (total 12 columns):  
#   Column                Non-Null Count  Dtype    
---  ---                  
0   price                 545 non-null   int64    
1   area                  545 non-null   int64    
2   bedrooms              545 non-null   int64    
3   bathrooms             545 non-null   int64    
4   stories               545 non-null   int64    
5   mainroad              545 non-null   object   
6   guestroom             545 non-null   object   
7   basement              545 non-null   object   
8   hotwaterheating       545 non-null   object   
9   airconditioning       545 non-null   object   
10  parking               545 non-null   int64    
11  furnishingstatus      545 non-null   object   
dtypes: int64(6), object(6)  
memory usage: 51.2+ KB
```

3. Perform Below Visualizations.

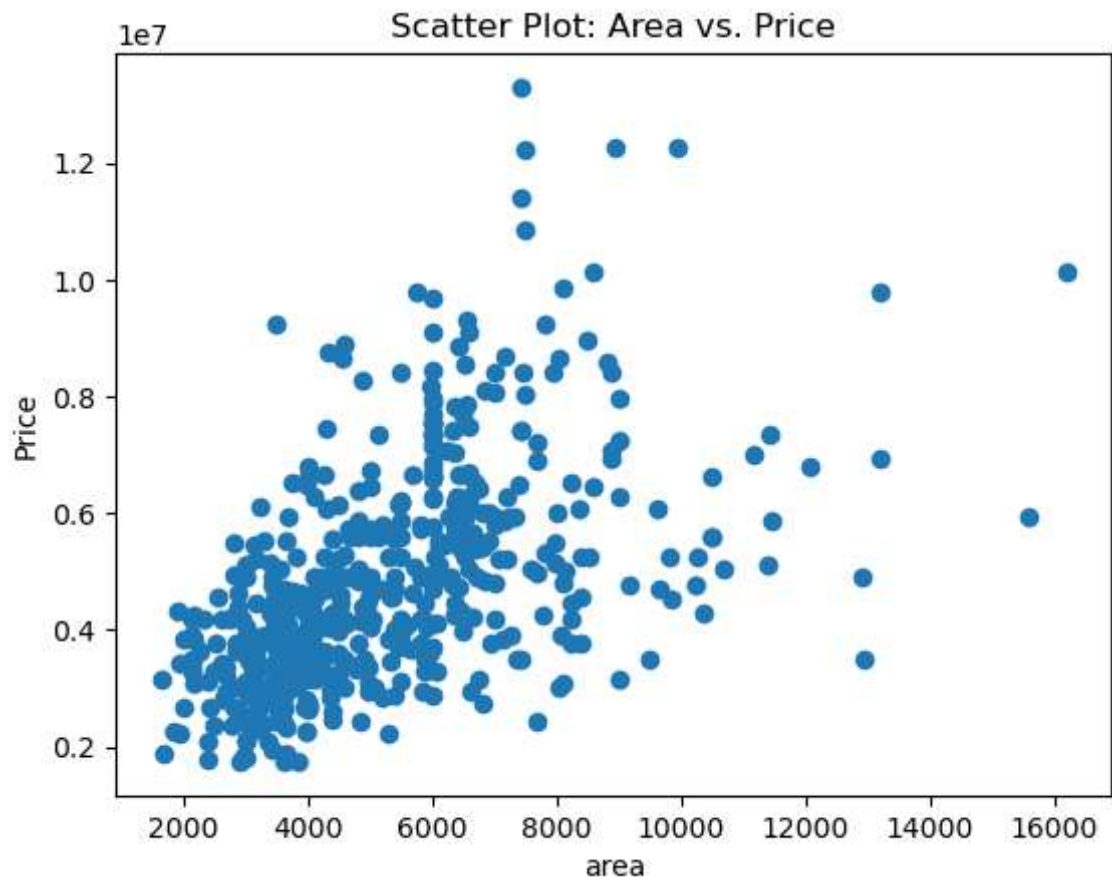
• Univariate Analysis

```
In [5]: plt.hist(df['price'], bins=20)
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.title('Distribution of Price')
plt.show()
```



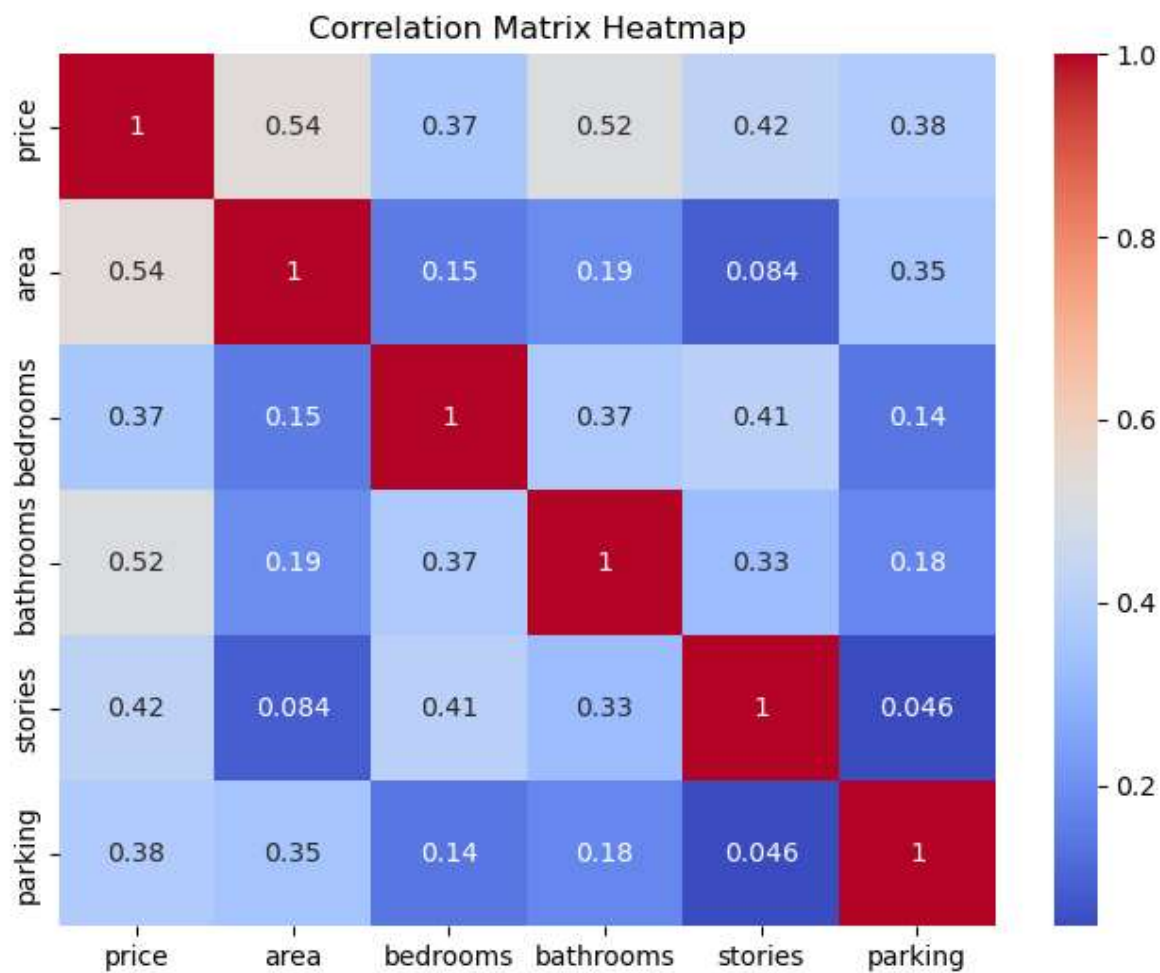
• Bi - Variate Analysis

```
In [6]: plt.scatter(df['area'], df['price'])  
plt.xlabel('area')  
plt.ylabel('Price')  
plt.title('Scatter Plot: Area vs. Price')  
plt.show()
```



• Multi - Variate Analysis

```
In [7]: corr_matrix = df.corr()  
plt.figure(figsize=(8, 6))  
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')  
plt.title('Correlation Matrix Heatmap')  
plt.show()
```



4. Perform descriptive statistics on the dataset.

In [8]: `df.describe()`

Out[8]:

	price	area	bedrooms	bathrooms	stories	parking
count	5.450000e+02	545.000000	545.000000	545.000000	545.000000	545.000000
mean	4.766729e+06	5150.541284	2.965138	1.286239	1.805505	0.693578
std	1.870440e+06	2170.141023	0.738064	0.502470	0.867492	0.861586
min	1.750000e+06	1650.000000	1.000000	1.000000	1.000000	0.000000
25%	3.430000e+06	3600.000000	2.000000	1.000000	1.000000	0.000000
50%	4.340000e+06	4600.000000	3.000000	1.000000	2.000000	0.000000
75%	5.740000e+06	6360.000000	3.000000	2.000000	2.000000	1.000000
max	1.330000e+07	16200.000000	6.000000	4.000000	4.000000	3.000000

5. Handle the Missing values.

In [9]: `df.isnull().sum()`

Out[9]:

price	0
area	0
bedrooms	0
bathrooms	0
stories	0
mainroad	0
guestroom	0
basement	0
hotwaterheating	0
airconditioning	0
parking	0
furnishingstatus	0
dtype: int64	

In [10]: `df.head()`

Out[10]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterhe
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	

6. Find the outliers and replace the outliers

```
In [11]: # numeric columns
numeric_cols = ['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'parking']

# Calculate the IQR for each column
Q1 = df[numeric_cols].quantile(0.25)
Q3 = df[numeric_cols].quantile(0.75)
IQR = Q3 - Q1

# Define the Lower and upper bounds for outliers
lower_bound = Q1 - (1.5 * IQR)
upper_bound = Q3 + (1.5 * IQR)

# Replace outliers with the median value
for col in numeric_cols:
    df.loc[(df[col] < lower_bound[col]) | (df[col] > upper_bound[col]), col] = df[col].median()

# Verify if outliers have been replaced
outliers_replaced = df[(df[numeric_cols] < lower_bound) | (df[numeric_cols] > upper_bound)]
print(outliers_replaced)
```

price	False
area	False
bedrooms	False
bathrooms	False
stories	False
mainroad	False
guestroom	False
basement	False
hotwaterheating	False
airconditioning	False
parking	False
furnishingstatus	False
dtype: bool	

7. Check for Categorical columns

```
In [12]: # Identify categorical columns
categorical_cols = df.select_dtypes(include='object').columns
print(categorical_cols)
df.head()
```

```
Index(['mainroad', 'guestroom', 'basement', 'hotwaterheating',
       'airconditioning', 'furnishingstatus'],
      dtype='object')
```

Out[12]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterhea
--	-------	------	----------	-----------	---------	----------	-----------	----------	-------------

0	4340000	7420	4	2	3	yes	no	no	
1	4340000	8960	4	1	2	yes	no	no	
2	4340000	9960	3	2	2	yes	no	yes	
3	4340000	7500	4	2	2	yes	no	yes	
4	4340000	7420	4	1	2	yes	yes	yes	



8. Split the data into dependent and independent variables.

```
In [13]: x = df.iloc[:,11] # Independent variables
y = df.iloc[:,11] # Dependent variable

# Display the independent variables (features)
print("Independent values:\n",x.head())

# Display the dependent variable (target)
print("\nDependent variable:\n",y.head())
```

Independent values:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement
0	4340000	7420	4	2	3	yes	no	no
1	4340000	8960	4	1	2	yes	no	no
2	4340000	9960	3	2	2	yes	no	yes
3	4340000	7500	4	2	2	yes	no	yes
4	4340000	7420	4	1	2	yes	yes	yes

	hotwaterheating	airconditioning	parking
0	no	yes	2
1	no	yes	0
2	no	no	2
3	no	yes	0
4	no	yes	2

Dependent variable:

0	furnished
1	furnished
2	semi-furnished
3	furnished
4	furnished

Name: furnishingstatus, dtype: object

perform encoding.

```
In [14]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
```

```
In [15]: ct=ColumnTransformer([('oh',OneHotEncoder()),[5,6,7,8,9]],remainder='passthro
```

```
In [16]: x=ct.fit_transform(x)
```

In [17]: `x.shape`

Out[17]: (545, 16)

In [18]: `x`

Out[18]: `array([[0., 1., 1., ..., 2., 3., 2.],
[0., 1., 1., ..., 1., 2., 0.],
[0., 1., 1., ..., 2., 2., 2.],
...,
[0., 1., 1., ..., 1., 1., 0.],
[1., 0., 1., ..., 1., 1., 0.],
[0., 1., 1., ..., 1., 2., 0.]])`

In [19]: `from sklearn.preprocessing import LabelEncoder`

In [20]: `le=LabelEncoder()`

In [21]: `y=le.fit_transform(y)`

In [22]: `y`

Out[22]: `array([0, 0, 1, 0, 0, 1, 1, 2, 0, 2, 0, 1, 1, 0, 1, 1, 2, 0, 0, 1, 1, 2,
0, 0, 0, 0, 1, 1, 2, 1, 2, 1, 0, 2, 0, 0, 0, 0, 2, 1, 0, 0, 2, 1,
0, 1, 0, 0, 2, 1, 2, 2, 0, 1, 1, 2, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0,
0, 2, 0, 0, 1, 2, 2, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 2, 1,
2, 0, 1, 0, 0, 1, 1, 1, 0, 1, 2, 2, 2, 2, 1, 0, 0, 2, 1, 2, 1, 1,
1, 2, 0, 0, 0, 1, 2, 0, 0, 1, 0, 1, 1, 0, 0, 1, 2, 2, 0, 1, 2, 1,
1, 2, 1, 2, 2, 1, 1, 0, 2, 1, 1, 2, 0, 0, 1, 1, 1, 1, 2, 1, 0, 1,
1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 2, 2, 1, 0, 1, 1, 1, 2, 2, 2, 1,
0, 1, 1, 1, 0, 1, 1, 0, 2, 0, 1, 2, 1, 0, 1, 0, 1, 1, 1, 1, 1,
1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 1, 0, 1, 1, 2, 2, 1, 0, 1, 2,
1, 1, 1, 0, 2, 1, 1, 0, 1, 1, 1, 1, 0, 2, 1, 1, 1, 1, 0, 1, 1,
0, 0, 1, 2, 1, 2, 1, 0, 1, 1, 1, 0, 1, 2, 0, 0, 1, 0, 0, 1, 1, 1,
1, 1, 2, 1, 1, 2, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 2, 0, 2, 0, 0, 1,
0, 2, 2, 1, 1, 2, 1, 2, 1, 1, 1, 2, 1, 0, 1, 1, 2, 0, 1, 1, 1, 1,
1, 1, 1, 1, 0, 1, 1, 1, 2, 2, 1, 1, 0, 1, 0, 1, 1, 1, 2, 1, 1, 0,
2, 0, 1, 1, 0, 0, 1, 2, 1, 0, 1, 1, 0, 1, 1, 1, 1, 2, 1, 1, 1, 1,
2, 0, 0, 0, 0, 0, 1, 2, 1, 1, 0, 1, 1, 0, 0, 2, 1, 0, 1, 1, 1, 0,
1, 2, 1, 2, 1, 2, 1, 1, 0, 0, 0, 2, 2, 1, 2, 1, 2, 1, 1, 2, 2, 2,
1, 0, 2, 1, 2, 2, 2, 2, 2, 2, 0, 2, 2, 2, 0, 2, 2, 2, 0, 2, 0,
2, 1, 2, 2, 2, 2, 1, 0, 0, 0, 2, 2, 2, 2, 0, 1, 1, 2, 0, 2, 1, 2,
2, 2, 2, 2, 0, 0, 2, 1, 1, 2, 1, 1, 1, 2, 2, 2, 1, 2, 2, 2, 2,
1, 2, 1, 2, 1, 1, 1, 0, 2, 2, 1, 2, 2, 1, 2, 2, 2, 2, 1, 0, 2, 1,
2, 2, 1, 2, 2, 1, 2, 0, 2, 0, 2, 2, 2, 2, 1, 2, 2, 2, 1, 1, 2, 2,
2, 2, 2, 0, 2, 2, 0, 2, 1, 2, 2, 2, 2, 2, 2, 0, 0, 2, 2, 2, 1,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 2, 0, 2])`

9. Scale the independent variables

```
In [23]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
```

```
In [24]: x=sc.fit_transform(x)

# Display the scaled independent variables
print(x)
```

```
[[-0.40562287  0.40562287  0.46531479 ...  1.47243614  2.2138449
  1.72906501]
 [-0.40562287  0.40562287  0.46531479 ... -0.57470084  0.56780742
 -0.79056181]
 [-0.40562287  0.40562287  0.46531479 ...  1.47243614  0.56780742
  1.72906501]
 ...
 [-0.40562287  0.40562287  0.46531479 ... -0.57470084 -1.07823005
 -0.79056181]
 [ 2.46534421 -2.46534421  0.46531479 ... -0.57470084 -1.07823005
 -0.79056181]
 [-0.40562287  0.40562287  0.46531479 ... -0.57470084  0.56780742
 -0.79056181]]
```

10. Split the data into training and testing

```
In [25]: from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,random_state=42)
```

```
In [26]: x_train
```

```
Out[26]: array([[ -0.40562287,  0.40562287,  0.46531479, ...,  1.47243614,
  0.56780742,  0.4692516 ],
 [-0.40562287,  0.40562287,  0.46531479, ...,  1.47243614,
 -1.07823005, -0.79056181],
 [-0.40562287,  0.40562287,  0.46531479, ..., -0.57470084,
 -1.07823005,  1.72906501],
 ...,
 [-0.40562287,  0.40562287,  0.46531479, ...,  1.47243614,
  2.2138449 ,  0.4692516 ],
 [-0.40562287,  0.40562287,  0.46531479, ..., -0.57470084,
 -1.07823005, -0.79056181],
 [-0.40562287,  0.40562287, -2.14908276, ...,  1.47243614,
  0.56780742,  0.4692516 ]])
```



```
In [31]: y_test
```

```
Out[31]: array([2, 0, 1, 1, 0, 2, 0, 0, 2, 0, 1, 2, 2, 1, 2, 2, 1, 2, 2, 1, 2, 1,
                2, 2, 1, 1, 2, 0, 0, 0, 1, 2, 0, 1, 2, 1, 1, 2, 2, 0, 1, 2, 0, 0,
                1, 2, 1, 1, 2, 2, 2, 2, 0, 1, 2, 0, 2, 2, 1, 2, 1, 2, 1, 2, 2, 1,
                0, 2, 2, 0, 2, 1, 0, 0, 2, 1, 1, 1, 0, 2, 0, 1, 0, 2, 2, 2, 2, 1,
                0, 1, 1, 0, 0, 1, 0, 0, 1, 2, 2, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0])
```

```
In [32]: y_test.shape
```

```
Out[32]: (109,)
```

Building a Regression Model

11. Build the Model

```
In [33]: from sklearn.svm import SVC
```

```
In [34]: model=SVC(kernel='rbf')
```

```
In [35]: model
```

```
Out[35]: SVC()
```

12. Train the Model

```
In [36]: # training the model
fit=model.fit(x_train,y_train)
```

13. Test the Model

```
In [37]: # test the model
pred=fit.predict(x_test)
```

```
In [38]: pred
```

```
Out[38]: array([1, 0, 2, 1, 2, 1, 1, 1, 2, 2, 1, 2, 2, 1, 1, 1, 2, 1, 1, 2, 1, 1,
                2, 1, 1, 1, 1, 1, 2, 1, 1, 0, 2, 1, 1, 1, 1, 2, 1, 1, 2, 0, 1,
                2, 1, 0, 1, 2, 0, 0, 2, 1, 1, 1, 2, 0, 2, 1, 2, 1, 1, 1, 0, 1, 1,
                1, 1, 1, 0, 1, 1, 2, 1, 2, 1, 1, 1, 1, 2, 0, 1, 1, 1, 2, 1, 2, 1,
                1, 1, 1, 0, 2, 1, 1, 0, 1, 0, 1, 1, 2, 0, 1, 1, 1, 1, 0, 0, 1])
```

```
In [39]: y_test
```

```
Out[39]: array([2, 0, 1, 1, 0, 2, 0, 0, 2, 0, 1, 2, 2, 1, 2, 2, 1, 2, 2, 1, 2, 1,
                2, 2, 1, 1, 2, 0, 0, 0, 1, 2, 0, 1, 2, 1, 1, 2, 2, 0, 1, 2, 0, 0,
                1, 2, 1, 1, 2, 2, 2, 2, 0, 1, 2, 0, 2, 2, 1, 2, 1, 2, 1, 2, 2, 1,
                0, 2, 2, 0, 2, 1, 0, 0, 2, 1, 1, 1, 0, 2, 0, 1, 0, 2, 2, 2, 2, 1,
                0, 1, 1, 0, 0, 1, 0, 0, 1, 2, 2, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0])
```

14. Measure the performance using Metrics.

```
In [40]: from sklearn.metrics import confusion_matrix, accuracy_score

print("confusion matrix:\n",confusion_matrix(y_test,pred))
print("accuracy:",accuracy_score(y_test,pred))
```

```
confusion matrix:
[[ 8 15  6]
 [ 3 30  6]
 [ 5 22 14]]
accuracy: 0.47706422018348627
```

```
In [ ]:
```